

华中科技大学

数据中心课程设计报告

院 系 计算机科学与技术学院

班 级 硕 2110 班

学 号 M202173849

姓 名 熊益

2022 年 01 月 07 日

一、系统搭建

1. 实验环境

操作系统：Windows 10 家庭中文版

处理器：11th Gen Intel(R) Core(TM) i7-11800H @ 2.30GHz 2.30 GHz

RAM：16GB

开发语言：Python

2. 服务端搭建

首先下载 Minio，并创建两个密钥 RootUser 和 RootPassword，获得对象存储资源访问的授权，在本次实验中设置为'admin'和'12345678'。

运行 Minio 时，将其工作在服务器模式，并将配置文件存入 Minio 所在的目录中，如图 1 所示。

```
D:\download\minio\minio_server>cd D:\download\minio\minio_server
D:\download\minio\minio_server>set MINIO_ROOT_USER=admin
D:\download\minio\minio_server>set MINIO_ROOT_PASSWORD=12345678
D:\download\minio\minio_server>set MINIO_PROMETHEUS_AUTH_TYPE=public
D:\download\minio\minio_server>minio.exe -C ./ server D:\download\minio\minio_server\data --console-address ":9090"
API: http://10.21.174.117:9000 http://127.0.0.1:9000
RootUser: admin
RootPass: 12345678

Console: http://10.21.174.117:9090 http://127.0.0.1:9090
RootUser: admin
RootPass: 12345678

Command-line: https://docs.min.io/docs/minio-client-quickstart-guide
$ mc.exe alias set myminio http://10.21.174.117:9000 admin 12345678
Documentation: https://docs.min.io
```

图 1 运行服务器

本次实验绑定的 IP 地址为 http://10.21.174.117:9000、http://127.0.0.1:9000，能通过其进行访问，如图 2 所示。

在网页上建一个 bucket：loadgen，如图 3 所示。

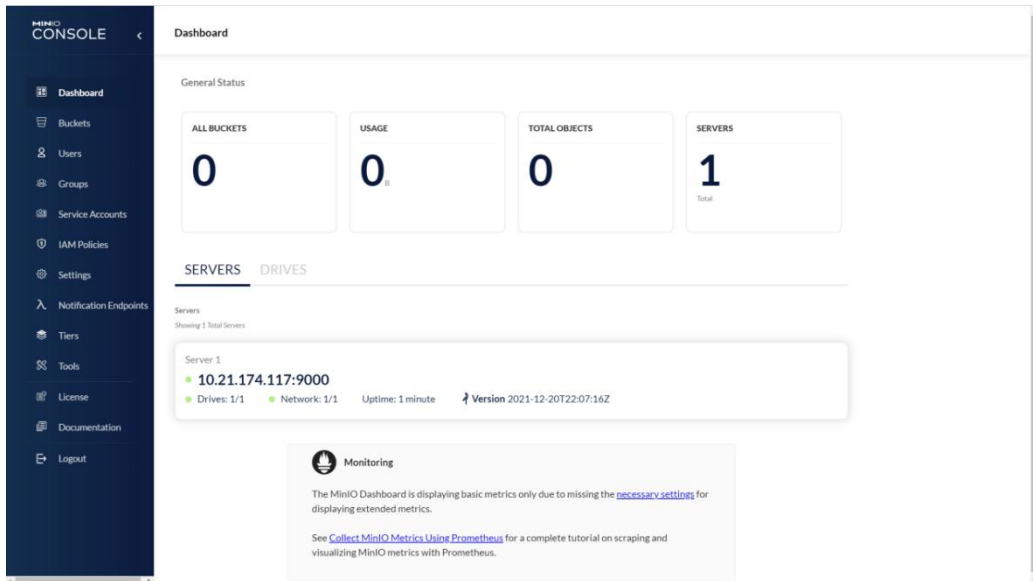


图 2 利用 IP 地址访问服务器

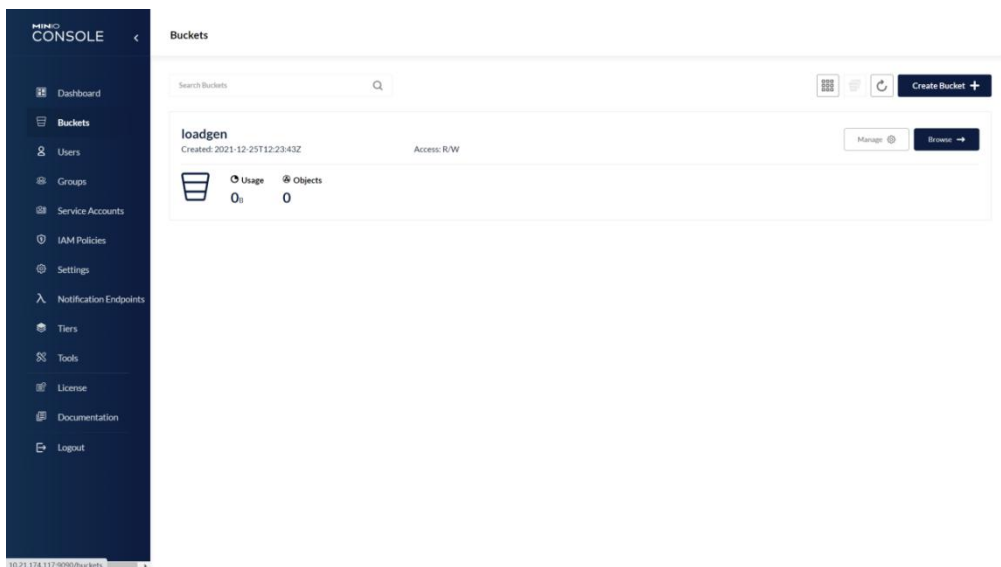


图 3 在网页上创建一个新的 bucket

二、性能观测

利用测评工具 S3 bench，在 bucket loadgen 中进行跑分，避免与用户数据重叠在一起，将 accessKey 和 accessSecret 修改为预设的 'admin' 和 '12345678'，endpoint 修改为本次实验绑定的 IP 地址。

初步设置为 8 个客户端，实验结果如图 4 所示。

从实验结果中观察到，8 个客户端的负载较轻，无法从数据中看出规律性，因此需要加重负载。

加客户端数量由 8 个上调至 32 个，实验结果如图 5 所示。

观察读延迟的数据，我们能发现，有 50% 的请求在 0.003s 内完成，但 90% 的请求在 0.006s 内完成，100% 的请求在 0.009s 内完成，即尾延迟现象，在写延迟中我们也能看到相同的现象。

```

D:\download\minio\minio_server\data>s3bench.exe ^
More? -accessKey=admin
More? -accessSecret=12345678 ^
More? -bucket=loadgen
More? -endpoint=http://127.0.0.1:9000 ^
More? -numClients=8
More? -numSamples=256 ^
More? -objectNamePrefix=loadgen ^
More? -objectSize=1024
Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:          loadgen
objectNamePrefix: loadgen
objectSize:      0.0010 MB
numClients:      8
numSamples:      256
verbose:         %!d(bool=false)

Generating in-memory sample data... Done (1.0017ms)

Running Write test...

Running Read test...

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:          loadgen
objectNamePrefix: loadgen
objectSize:      0.0010 MB
numClients:      8
numSamples:      256
verbose:         %!d(bool=false)

Results Summary for Write Operation(s)
Total Transferred: 0.250 MB
Total Throughput:  0.29 MB/s
Total Duration:    0.866 s
Number of Errors:  0
-----
Write times Max:    0.084 s
Write times 99th %ile: 0.081 s
Write times 90th %ile: 0.051 s
Write times 75th %ile: 0.039 s
Write times 50th %ile: 0.020 s
Write times 25th %ile: 0.009 s
Write times Min:    0.004 s

Results Summary for Read Operation(s)
Total Transferred: 0.250 MB
Total Throughput:  10.36 MB/s
Total Duration:    0.024 s
Number of Errors:  0
-----
Read times Max:     0.003 s
Read times 99th %ile: 0.003 s
Read times 90th %ile: 0.001 s
Read times 75th %ile: 0.001 s
Read times 50th %ile: 0.001 s
Read times 25th %ile: 0.001 s
Read times Min:     0.000 s

Cleaning up 256 objects...
Deleting a batch of 256 objects in range {0, 255}... Succeeded
Successfully deleted 256/256 objects in 107.3407ms

```

图 4 设置 8 个客户端

```
D:\download\minio\minio_server\data>s3bench.exe ^
More? -accessKey=admin ^
More? -accessSecret=12345678 ^
More? -bucket=loadgen ^
More? -endpoint=http://127.0.0.1:9000 ^
More? -numClients=32 ^
More? -numSamples=256 ^
More? -objectNamePrefix=loadgen ^
More? -objectSize=1024
Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.0010 MB
numClients:       32
numSamples:       256
verbose:          %!d(bool=false)

Generating in-memory sample data... Done (0s)

Running Write test...

Running Read test...

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.0010 MB
numClients:       32
numSamples:       256
verbose:          %!d(bool=false)

Results Summary for Write Operation(s)
Total Transferred: 0.250 MB
Total Throughput:  0.13 MB/s
Total Duration:    1.902 s
Number of Errors:  0
-----
Write times Max:      0.438 s
Write times 99th %ile: 0.383 s
Write times 90th %ile: 0.290 s
Write times 75th %ile: 0.261 s
Write times 50th %ile: 0.220 s
Write times 25th %ile: 0.197 s
Write times Min:      0.010 s

Results Summary for Read Operation(s)
Total Transferred: 0.250 MB
Total Throughput:  10.60 MB/s
Total Duration:    0.024 s
Number of Errors:  0
-----
Read times Max:       0.009 s
Read times 99th %ile: 0.006 s
Read times 90th %ile: 0.004 s
Read times 75th %ile: 0.004 s
Read times 50th %ile: 0.003 s
Read times 25th %ile: 0.002 s
Read times Min:       0.001 s

Cleaning up 256 objects...
Deleting a batch of 256 objects in range {0, 255}... Succeeded
Successfully deleted 256/256 objects in 106.659ms
```

图 5 上调至 32 个客户端

三、观测尾延迟

安装 Anaconda, 并将其配置到 VS Code 环境中, 在运行 latency-collect 代码前, 需要在 Anaconda 中使用 `pip install` 指令安装 `botocore`、`tqdm`、`throttle` 包, 如图 6 所示。

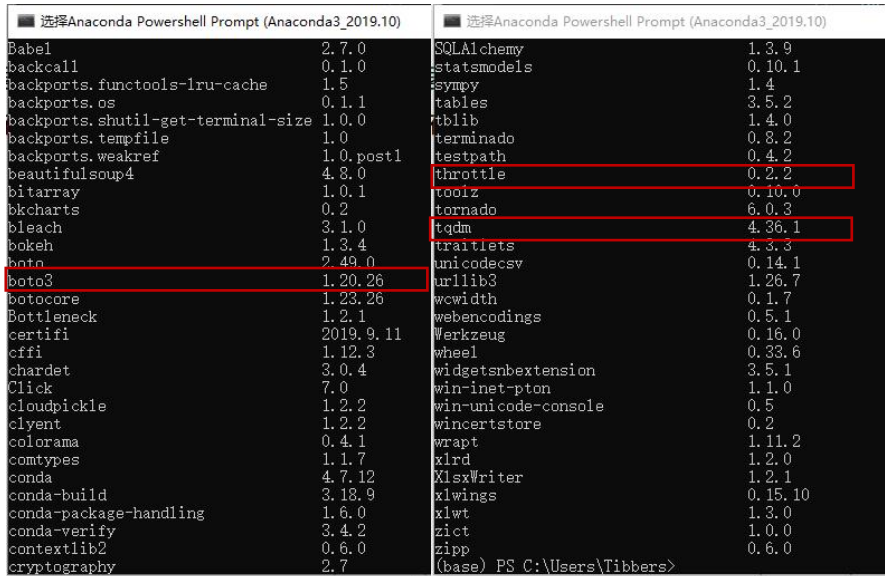


图 6 Anaconda 安装包

在 python 环境中建立与服务器的对话。

```
# 准备密钥
aws_access_key_id = 'admin'
aws_secret_access_key = '12345678'

# 本地s3服务地址
local_s3 = 'http://127.0.0.1:9000'

# 建立会话
session = Session(aws_access_key_id=aws_access_key_id, aws_secret_access_key=aws_secret_access_key)

# 连接到服务
s3 = session.resource('s3', endpoint_url=local_s3)
```

新建一个实验用的 bucket, 并查看所有 bucket

```
21
22 #新建一个实验用 bucket (注意: "bucket name" 中不能有下划线)
23 bucket_name = 'test100objs'
24 if s3.Bucket(bucket_name) not in s3.buckets.all():
25     s3.create_bucket(Bucket=bucket_name)
26
27 #查看所有pocket
28 for bucket in s3.buckets.all():
29     print('bucket name:%s' % bucket.name)
30
```

问题 输出 调试控制台 终端 Code

```
[Running] python -u "d:\download\vscode\workplace\test\test.py"
bucket name:loadgen
bucket name:test100objs

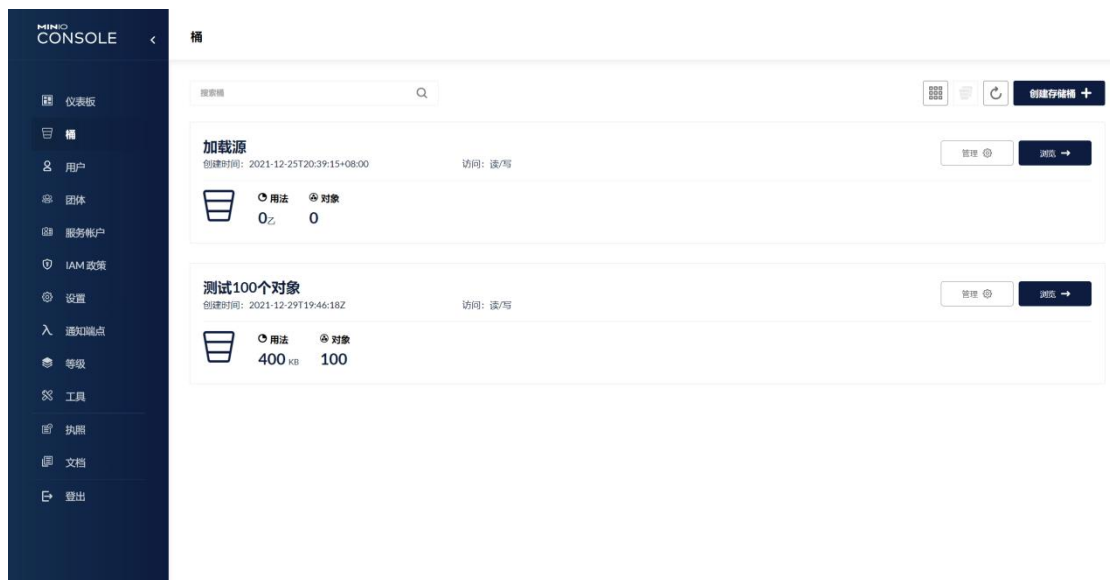
[Done] exited with code=0 in 0.329 seconds
```


准备负载，并按照与设 IAT 发起请求

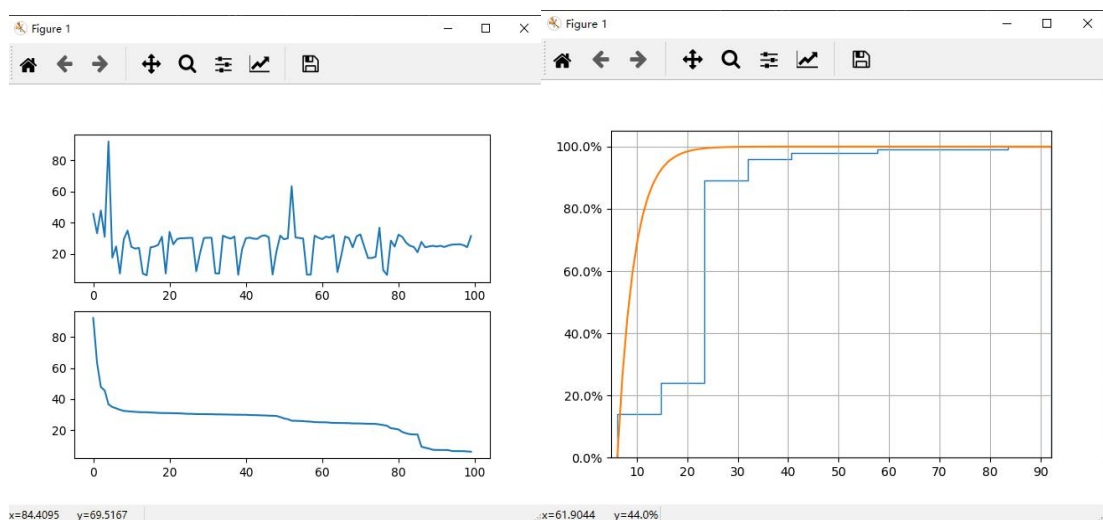
```
[Running] python -u "d:\download\vscode\workplace\test\test.py"
bucket name:loadgen
bucket name:test100objs

Accessing S3: 0%|          | 0/100 [00:00<?, ?it/s]
Accessing S3: 1%|          | 1/100 [00:00<00:35, 2.81it/s]
Accessing S3: 16%|██████    | 16/100 [00:00<00:21, 3.97it/s]
Accessing S3: 20%|██████    | 20/100 [00:00<00:14, 5.40it/s]
Accessing S3: 24%|██████    | 24/100 [00:00<00:10, 7.19it/s]
Accessing S3: 28%|██████    | 28/100 [00:00<00:07, 9.37it/s]
Accessing S3: 35%|██████    | 35/100 [00:00<00:05, 12.49it/s]
Accessing S3: 40%|██████    | 40/100 [00:01<00:03, 15.29it/s]
Accessing S3: 44%|██████    | 44/100 [00:01<00:03, 18.04it/s]
Accessing S3: 48%|██████    | 48/100 [00:01<00:02, 20.80it/s]
Accessing S3: 52%|██████    | 52/100 [00:01<00:01, 24.24it/s]
Accessing S3: 56%|██████    | 56/100 [00:01<00:01, 26.48it/s]
Accessing S3: 61%|██████    | 61/100 [00:01<00:01, 29.59it/s]
Accessing S3: 66%|██████    | 66/100 [00:01<00:01, 32.20it/s]
Accessing S3: 70%|██████    | 70/100 [00:01<00:00, 32.16it/s]
Accessing S3: 74%|██████    | 74/100 [00:02<00:00, 31.87it/s]
Accessing S3: 78%|██████    | 78/100 [00:02<00:00, 32.11it/s]
Accessing S3: 82%|██████    | 82/100 [00:02<00:00, 32.27it/s]
Accessing S3: 88%|██████    | 88/100 [00:02<00:00, 35.60it/s]
Accessing S3: 92%|██████    | 92/100 [00:02<00:00, 34.70it/s]
Accessing S3: 97%|██████    | 97/100 [00:02<00:00, 36.23it/s]
Accessing S3: 100%|██████   | 100/100 [00:02<00:00, 35.76it/s]

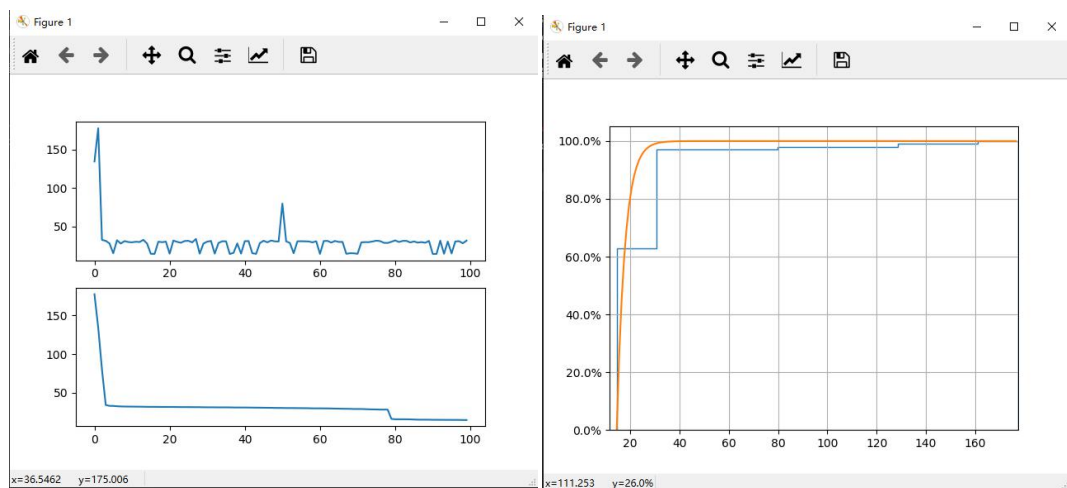
[Done] exited with code=0 in 3.156 seconds
```



当设置文件设置为 4k，并将 latency 存入文件并，并运行 latency-plot 的代码，将数据绘制为图表，下图为是实验结果：



将文件大小上调至 32k 后的实验结果为：



左边图表中的上方图表将请求依次打印出来，我们可以看到前期的波动有剧烈的波动，中间也存在较大波动的地方，这些地方的延迟超过了远平均延迟之上。下方表格将其排序，能看出只有少部分的延迟较高，大部分地方的延迟为短时间延迟。

右侧的表格可以看出多达 65% 的请求能够早 10ms 以内完成，但 90% 的请求要在 30ms 以内完成，若请求全部完成，则需要 160ms，这也就是实验中的长尾延迟现象。