



# 华中科技大学

## 数据中心技术实验报告

姓 名： 陈倩  
学 院： 计算机科学与技术学院  
专 业： 电子信息  
班 级： 2110  
学 号： M202173864  
指导教师： 施展、童薇

分数	
教师签名	

2021 年 12 月 26 日

# 目 录

<b>1</b>	<b>实验目的及实验内容 .....</b>	<b>1</b>
<b>2</b>	<b>实验一：系统搭建 .....</b>	<b>2</b>
2.1	实验环境.....	2
2.2	基础环境.....	2
2.3	服务端搭建.....	2
<b>3</b>	<b>实验二：性能观测 .....</b>	<b>3</b>
3.1	评测工具.....	3
3.2	标准测试.....	3
3.2.1	对象尺寸对性能的影响.....	3
3.2.2	客户端数量对性能的影响 .....	6
<b>4</b>	<b>实验三：尾延迟测试 .....</b>	<b>12</b>
4.1	实验环境.....	12
4.2	延迟收集.....	13
4.3	尾延迟展示.....	13

# 1 实验目的及实验内容

对象存储 (Object Storage Service, OSS), 也叫基于对象的存储, 是一种解决和处理离散单元的方法, 可提供基于分布式系统之上的对象形式的数据存储服务。对象存储和我们经常接触到的块和文件系统等存储形态不同, 它提供 RESTful API 数据读写接口及丰富的 SDK 接口, 并且常以网络服务的形式提供数据的访问。作为新数据时代的新存储形态, 相比 SAN 和 NAS, 对象存储直接提供 API 给应用使用, 采用扁平化的结构管理所有桶 (Bucket) 和对象。每个桶和对象都有一个全局唯一的 ID, 根据 ID 可快速实现对象的查找和数据的访问。对象存储支持基于策略的自动化管理机制, 使得每个应用可根据业务需要动态地控制每个桶的数据冗余策略、数据访问权限控制及数据生命周期管理。这些优势使得对象存储具备极致的扩展和极易的数据管理。本实验是对象存储的一个入门实践。

实验内容包括:

1. 熟悉性能指标: 吞吐率、带宽、延迟
2. 分析不同负载下的指标、延迟的分布
3. 观测尾延迟现象
4. 尝试对冲请求方案

## 2 实验一：系统搭建

### 2.1 实验环境

VMware 版本：VMware Workstation 16 Player 16.1.2

Linux 版本：5.11.0-37-generic

### 2.2 基础环境

Python 版本：3.8.10

Go 版本：go1.10.3 linux/amd64jupyter

Jupyter notebook 版本：6.4.6

### 2.3 服务端搭建

首先将脚本（[https://gitee.com/shi\\_zhan/obs-tutorial](https://gitee.com/shi_zhan/obs-tutorial)）克隆到本地，从 minio 官网下载 Linux 版本的 minio 到 obs-tutorial 文件夹，然后进入 obs-tutorial 文件，在命令行键入 ./run-minio.sh 启动 minio。run-minio.sh 指定 minio 用户名和密码分别为 hust 和 hust\_obs，端口为 9090。启动 minio 后，在浏览器内输入 127.0.0.1: 9090 打开 minio 登录界面，输入用户名和密码后，进入 minio 主页。

```
#!/bin/bash

# Use environment variables MINIO_ROOT_USER & MINIO_ROOT_PASSWORD to set keys, for later use in clients.
export MINIO_ROOT_USER=hust
export MINIO_ROOT_PASSWORD=hust_obs

# Export metrics
export MINIO_PROMETHEUS_AUTH_TYPE="public"

# Use "-C" flag to store configuration file in local directory "./".
# Use server command to start object storage server with "./root" as root directory, in which holds all buckets and objects.
./minio -C ./ server ./root --console-address ":9090"

# Run above task in one command line.
# export MINIO_ROOT_USER=hust && export MINIO_ROOT_PASSWORD=hust_obs && ./minio -C ./ server ./root
```

图 2-1 run-minio.sh

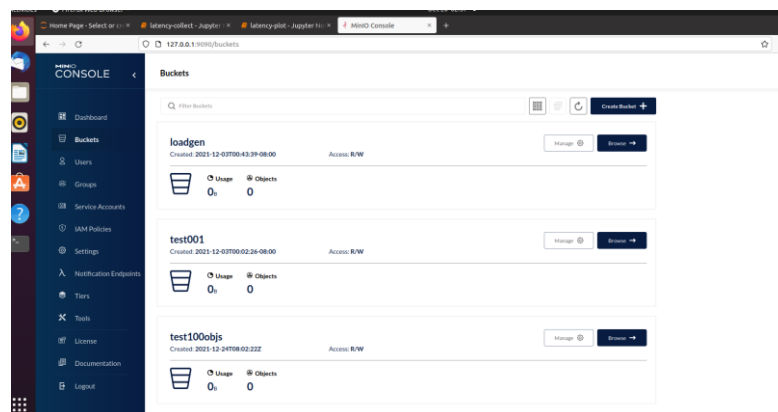


图 2-2 minio 主页

## 3 实验二：性能观测

### 3.1 评测工具

本实验选择 S3 Bench 作为评测工具。在安装 S3 Bench 之前，需要先安装 Go，首先去官网下载最新版 Golang 的安装包，下载完毕后，解压到/usr/local 下。然后编辑/etc/profile 文件配置环境变量。

```
export GOROOT=/usr/local/go
export PATH=$PATH:$GOROOT/bin

export GOPATH=$HOME/go
export PATH=$PATH:$GOPATH/bin

export PATH=$PATH:~/local/bin
```

图 3-1 Golang 环境变量

配置好 Go 后，在命令行键入 `go get -u github.com/igneous-systems/s3bench` 下载 s3bench，s3bench 的默认安装路径为 `~/go/bin/s3bench`。安装完成后，将 obs-tutorial 下的 run-s3bench.sh 的 s3bench 路径修改为 `~/go/bin/s3bench`。

```
s3bench=~/go/bin/s3bench

# if [ -n "$GOPATH" ]; then
#     s3bench=/home/monica/go/bin/s3bench
# fi
```

图 3-2 s3bench 路径

### 3.2 标准测试

指标：吞吐率 Throughput、延迟 Latency、对象尺寸 ObjectSize、客户端数量、服务器数量等。

```
# -accessKey      Access Key
# -accessSecret   Secret Key
# -bucket=loadgen Bucket for holding all test objects.
# -endpoint=http://127.0.0.1:9000 Endpoint URL of object storage service being tested.
# -numClients=8   Simulate 8 clients running concurrently.
# -numSamples=256 Test with 256 objects.
# -objectNamePrefix=loadgen Name prefix of test objects.
# -objectSize=1024 Size of test objects.
# -verbose        Print latency for every request.
```

图 3-3 s3bench 各项参数涵义

#### 3.2.1 对象尺寸对性能的影响

通过修改 run-s3bench.sh 的 ObjectSize 参数来改变对象尺寸，实验选择的对象尺寸为 1024\*32、1024\*320 和 1024\*3200。

```
Generating in-memory sample data... Done (277.686µs)
```

```
Running Write test...
```

```
Running Read test...
```

```
Test parameters
```

```
endpoint(s):      [http://127.0.0.1:9000]  
bucket:           loadgen  
objectNamePrefix: loadgen  
objectSize:       0.0312 MB  
numClients:       8  
numSamples:       256  
verbose:          %!d(bool=false)
```

```
Results Summary for Write Operation(s)
```

```
Total Transferred: 8.000 MB  
Total Throughput:  18.19 MB/s  
Total Duration:    0.440 s  
Number of Errors:  0
```

```
-----  
Write times Max:      0.128 s  
Write times 99th %ile: 0.087 s  
Write times 90th %ile: 0.025 s  
Write times 75th %ile: 0.018 s  
Write times 50th %ile: 0.010 s  
Write times 25th %ile: 0.005 s  
Write times Min:      0.002 s
```

图 3-4 ObjectSize=(1024\*32)(a)

```
Results Summary for Read Operation(s)
```

```
Total Transferred: 8.000 MB  
Total Throughput:  45.35 MB/s  
Total Duration:    0.176 s  
Number of Errors:  0
```

```
-----  
Read times Max:       0.060 s  
Read times 99th %ile: 0.048 s  
Read times 90th %ile: 0.013 s  
Read times 75th %ile: 0.005 s  
Read times 50th %ile: 0.002 s  
Read times 25th %ile: 0.001 s  
Read times Min:       0.001 s
```

```
Cleaning up 256 objects...
```

```
Deleting a batch of 256 objects in range {0, 255}... Succeeded  
Successfully deleted 256/256 objects in 149.7574ms  
monica@ubuntu:~/Desktop/data-center/obs-tutorial$
```

图 3-5 ObjectSize=(1024\*32)(b)

```

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.3125 MB
numClients:       8
numSamples:       256
verbose:          %!d(bool=false)

Generating in-memory sample data... Done (2.596528ms)

Running Write test...

Running Read test...

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.3125 MB
numClients:       8
numSamples:       256
verbose:          %!d(bool=false)

Results Summary for Write Operation(s)
Total Transferred: 80.000 MB
Total Throughput:  109.50 MB/s
Total Duration:    0.731 s
Number of Errors:  0
-----
Write times Max:      0.109 s
Write times 99th %ile: 0.090 s
Write times 90th %ile: 0.045 s
Write times 75th %ile: 0.026 s
Write times 50th %ile: 0.017 s
Write times 25th %ile: 0.011 s
Write times Min:      0.003 s

```

图 3-6 ObjectSize=(1024\*320)(a)

```

Results Summary for Read Operation(s)
Total Transferred: 80.000 MB
Total Throughput:  355.48 MB/s
Total Duration:    0.225 s
Number of Errors:  0
-----
Read times Max:       0.060 s
Read times 99th %ile: 0.040 s
Read times 90th %ile: 0.021 s
Read times 75th %ile: 0.009 s
Read times 50th %ile: 0.003 s
Read times 25th %ile: 0.001 s
Read times Min:       0.001 s

Cleaning up 256 objects...
Deleting a batch of 256 objects in range {0, 255}... Succeeded
Successfully deleted 256/256 objects in 103.906741ms

```

图 3-7 ObjectSize=(1024\*320)(b)

```

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       3.1250 MB
numClients:       8
numSamples:       256
verbose:          %!d(bool=false)

Generating in-memory sample data... Done (37.450165ms)

Running Write test...

Running Read test...

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       3.1250 MB
numClients:       8
numSamples:       256
verbose:          %!d(bool=false)

Results Summary for Write Operation(s)
Total Transferred: 800.000 MB
Total Throughput:  81.24 MB/s
Total Duration:    9.847 s
Number of Errors:  0
-----
Write times Max:      1.923 s
Write times 99th %ile: 1.862 s
Write times 90th %ile: 0.506 s
Write times 75th %ile: 0.367 s
Write times 50th %ile: 0.205 s
Write times 25th %ile: 0.117 s
Write times Min:      0.035 s

```

图 3-8 ObjectSize=(1024\*3200)(a)

```

Results Summary for Read Operation(s)
Total Transferred: 800.000 MB
Total Throughput:  249.36 MB/s
Total Duration:    3.208 s
Number of Errors:  0
-----
Read times Max:       0.717 s
Read times 99th %ile: 0.704 s
Read times 90th %ile: 0.208 s
Read times 75th %ile: 0.109 s
Read times 50th %ile: 0.061 s
Read times 25th %ile: 0.041 s
Read times Min:       0.005 s

Cleaning up 256 objects...
Deleting a batch of 256 objects in range {0, 255}... Succeeded
Successfully deleted 256/256 objects in 1.602345079s
monica@ubuntu:~/Desktop/data-center/obs-tutorial$

```

图 3-9 ObjectSize=(1024\*3200)(b)

通过实验对比发现，对象尺寸越大，吞吐量、延迟也会增大（也受网络情况影响），同时，长尾效应也会更加严重。

### 3.2.2 客户端数量对性能的影响

通过修改 run-s3bench.sh 的 numClients 参数来改变客户端数，实验选择的客户端数为 8、32、64、128 和 256。



```

Test parameters
endpoint(s): [http://127.0.0.1:9000]
bucket: loadgen
objectNamePrefix: loadgen
objectSize: 0.0312 MB
numClients: 8
numSamples: 256
verbose: %!d(bool=false)

Generating in-memory sample data... Done (590.069µs)

Running Write test...

Running Read test...

Test parameters
endpoint(s): [http://127.0.0.1:9000]
bucket: loadgen
objectNamePrefix: loadgen
objectSize: 0.0312 MB
numClients: 8
numSamples: 256
verbose: %!d(bool=false)

Results Summary for Write Operation(s)
Total Transferred: 8.000 MB
Total Throughput: 22.90 MB/s
Total Duration: 0.349 s
Number of Errors: 0
-----
Write times Max: 0.041 s
Write times 99th %ile: 0.029 s
Write times 90th %ile: 0.020 s
Write times 75th %ile: 0.015 s
Write times 50th %ile: 0.010 s
Write times 25th %ile: 0.006 s
Write times Min: 0.001 s

```

图 3-10 numClients=8(a)

```

Results Summary for Read Operation(s)
Total Transferred: 8.000 MB
Total Throughput: 60.41 MB/s
Total Duration: 0.132 s
Number of Errors: 0
-----
Read times Max: 0.039 s
Read times 99th %ile: 0.034 s
Read times 90th %ile: 0.011 s
Read times 75th %ile: 0.005 s
Read times 50th %ile: 0.002 s
Read times 25th %ile: 0.001 s
Read times Min: 0.000 s

```

图 3-11 numClients=8(b)

```

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.0312 MB
numClients:       32
numSamples:       256
verbose:          %!d(bool=false)

Generating in-memory sample data... Done (314.025µs)

Running Write test...

Running Read test...

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.0312 MB
numClients:       32
numSamples:       256
verbose:          %!d(bool=false)

Results Summary for Write Operation(s)
Total Transferred: 8.000 MB
Total Throughput:  25.65 MB/s
Total Duration:    0.312 s
Number of Errors:  0
-----
Write times Max:      0.124 s
Write times 99th %ile: 0.090 s
Write times 90th %ile: 0.066 s
Write times 75th %ile: 0.048 s
Write times 50th %ile: 0.035 s
Write times 25th %ile: 0.023 s
Write times Min:      0.002 s

```

图 3-12 numClients=32(a)

```

Results Summary for Read Operation(s)
Total Transferred: 8.000 MB
Total Throughput:  43.58 MB/s
Total Duration:    0.184 s
Number of Errors:  0
-----
Read times Max:      0.139 s
Read times 99th %ile: 0.135 s
Read times 90th %ile: 0.068 s
Read times 75th %ile: 0.038 s
Read times 50th %ile: 0.004 s
Read times 25th %ile: 0.001 s
Read times Min:      0.001 s

```

图 3-13 numClients=32(b)

```

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.0312 MB
numClients:       64
numSamples:       256
verbose:          %!d(bool=false)

Generating in-memory sample data... Done (219.623µs)

Running Write test...

Running Read test...

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.0312 MB
numClients:       64
numSamples:       256
verbose:          %!d(bool=false)

Results Summary for Write Operation(s)
Total Transferred: 8.000 MB
Total Throughput:  22.46 MB/s
Total Duration:    0.356 s
Number of Errors:  0
-----
Write times Max:      0.160 s
Write times 99th %ile: 0.144 s
Write times 90th %ile: 0.109 s
Write times 75th %ile: 0.093 s
Write times 50th %ile: 0.080 s
Write times 25th %ile: 0.063 s
Write times Min:      0.009 s

```

图 3-14 numClients=64(a)

```

Results Summary for Read Operation(s)
Total Transferred: 8.000 MB
Total Throughput:  56.32 MB/s
Total Duration:    0.142 s
Number of Errors:  0
-----
Read times Max:       0.088 s
Read times 99th %ile: 0.078 s
Read times 90th %ile: 0.054 s
Read times 75th %ile: 0.045 s
Read times 50th %ile: 0.030 s
Read times 25th %ile: 0.019 s
Read times Min:       0.001 s

```

图 3-15 numClients=64(b)

```

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.0312 MB
numClients:       128
numSamples:       256
verbose:          %!d(bool=false)

```

```

Results Summary for Write Operation(s)
Total Transferred: 8.000 MB
Total Throughput:  16.18 MB/s
Total Duration:    0.494 s
Number of Errors:  0

```

```

-----
Write times Max:      0.406 s
Write times 99th %ile: 0.393 s
Write times 90th %ile: 0.350 s
Write times 75th %ile: 0.305 s
Write times 50th %ile: 0.207 s
Write times 25th %ile: 0.104 s
Write times Min:      0.021 s

```

```

Results Summary for Read Operation(s)
Total Transferred: 8.000 MB
Total Throughput:  61.44 MB/s
Total Duration:    0.130 s
Number of Errors:  0

```

```

-----
Read times Max:       0.111 s
Read times 99th %ile: 0.110 s
Read times 90th %ile: 0.104 s
Read times 75th %ile: 0.094 s
Read times 50th %ile: 0.054 s
Read times 25th %ile: 0.013 s
Read times Min:       0.001 s

```

图 3-16 numClients=128

```

Test parameters
endpoint(s):      [http://127.0.0.1:9000]
bucket:           loadgen
objectNamePrefix: loadgen
objectSize:       0.0312 MB
numClients:       256
numSamples:       256
verbose:          %!d(bool=false)

```

```

Results Summary for Write Operation(s)
Total Transferred: 8.000 MB
Total Throughput:  24.53 MB/s
Total Duration:    0.326 s
Number of Errors:  0

```

```

-----
Write times Max:      0.270 s
Write times 99th %ile: 0.249 s
Write times 90th %ile: 0.225 s
Write times 75th %ile: 0.199 s
Write times 50th %ile: 0.141 s
Write times 25th %ile: 0.108 s
Write times Min:      0.066 s

```

```

Results Summary for Read Operation(s)
Total Transferred: 8.000 MB
Total Throughput:  40.35 MB/s
Total Duration:    0.198 s
Number of Errors:  0

```

```

-----
Read times Max:       0.186 s
Read times 99th %ile: 0.185 s
Read times 90th %ile: 0.174 s
Read times 75th %ile: 0.154 s
Read times 50th %ile: 0.128 s
Read times 25th %ile: 0.113 s
Read times Min:       0.032 s

```

图 3-17 numClients=256

通过实验对比发现，在 256 以内，客户端数越大，吞吐量、总延迟的波动情况不是特别明显，但是，随着客户端的增加，长尾效应会更加严重。

## 4 实验三：尾延迟测试

### 4.1 实验环境

Python 版本：3.8.10

Jupyter notebook 版本：6.4.6

首先在命令行键入 `pip3 install jupyter` 下载 python3.8 对应版本的 jupyter，然后把 jupyter 的安装路径添加到系统环境变量/etc/profile 中。

```
export PATH=$PATH:~/local/bin
```

图 4-1 jupyter 路径

在命令行键入 `pip3 install ipython` 安装 ipython 模块，终端输入 `ipython` 进入 ipython 界面，为 jupyter 配置密码，输入 `from notebook.auth import passwd`，然后输入密码，将生成的 SHA1 密钥复制下来。然后，在终端输入命令：`jupyter-notebook --generate-config --allow-root` 生成配置文件，默认地址为 `./jupyter/jupyter_notebook_config.py`，使用 `vim` 编辑配置文件。

```
#Nginx访问时会出现跨域访问，需要在这里允许
c.NotebookApp.allow_origin = '*'

# 禁止随意修改密码
c.NotebookApp.allow_password_change = False

# 是否允许远程访问
c.NotebookApp.allow_remote_access = True

# IP
c.NotebookApp.ip = '0.0.0.0'

# 端口
c.NotebookApp.port = 9820

# 工作目录
c.NotebookApp.notebook_dir = 'jupyter/'

# 启动Jupyter Notebook之后是否打开浏览器
c.NotebookApp.open_browser = False

# 客户端打开Jupyter Notebook的密码哈希值
c.NotebookApp.password = 'argon2:$argon2id$v=19$m=10240,t=10,p=8$H1qLceYjd8voWf8iLY5dQA5je+oH+zrJcyEUSAUEK/Hj6nSggIkdTfVl3aHP3Dqe4'
```

图 4-2 jupyter notebook 配置文件

在终端输入命令：`jupyter-notebook --allow-root` 启动 jupyter notebook,然后在浏览器中打开 `127.0.0.1: 9820` 进入 jupyter notebook 首页。

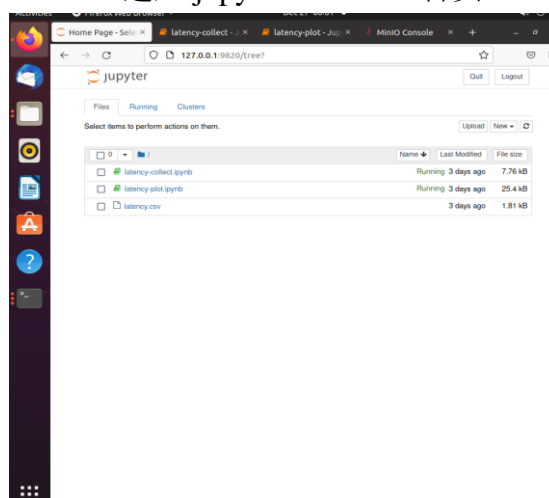


图 4-3 jupyter notebook 首页

## 4.2 延迟收集

将 obs-tutorial 文件下的 latency-collect.py 上传到 jupyter notebook 上，依次运行并获得 100 个 4Kb 对象的写延迟。




图 4-4 收集到的 latency.csv

## 4.3 尾延迟展示

将 obs-tutorial 文件下的 latency-plot.py 上传到 jupyter notebook 上，依次运行得到延迟分布情况。




图 4-5 延迟分布情况（上图为无序状态，下图为有序状态）




图 4-6 百分位延迟